

First test with XOR.

Source data: *XOR.json*.

I built the network using the standard setup for XOR; 2 inputs, 2 hidden nodes and 1 output. I adjusted the learning rate to minimise predictive errors, finding a learning rate of 0.2 to work best.

Experiments with weight adjustments

Initially built the MLP using stochastic gradient descent; but had issues while testing with XOR where about half of all tests would get stuck in local minima. This resulted in an output like;

Outcome	Expected
[[0.05774859]]	0
[[0.57647176]]	0
[[0.96448377]]	1
[[0.57119892]]	1

And the other half perfect:

Outcome	Expected
[[0.01585943]]	0
[[0.01289982]]	0
[[0.98725898]]	1
[[0.987264]]	1

With an error readout as one might expect; decreasing for all data in the second instance and hovering about 0.5 for two of four samples in the 1st instance. I think I would need to have a decreasing learning rate for STG to work effectively.

I then tried Gradient Descent where I updated 25% of the time and played with this figure until I was getting a correct reading out 75% of the time; at 10% of the time. By correct I mean output is close to expected.

Conclusion

Final iteration was running with weights updating about 15 of 100 of every iteration through the dataset and a learning rate of 0.2. This seemed to produce the most consistent correct results.

Testing with Generated Vectors

Source data: *vector.json*.

I imported the model class from NeuralNetwork.py and created an instance with 4 input nodes, 5 hidden nodes and one output node and learning rate of 0.2. I ran through the data for 1000 epochs.

The prediction, expected result and error for each prediction were printed to the file *Vectors_Predict_timestamp.txt*.

The final prediction and error are given below:

Results

Prediction	Expected	Error
0.91558071	0.9025726658005266	-0.01300805

The errors enjoyed a downward trajectory over the course of each iteration. The first datapoint of the first epoch had an error of -1.139886 which was close to 0 by the 1000th epoch, final error given was 0.07695601.

Conclusion

The model performed very strongly. I was obliged to reduce the number of iterations somewhat to produce a result in a reasonable time but at 100 epochs the results remained consistently good with this model.